

Cassandra Day

Application Development



Source Code Inside

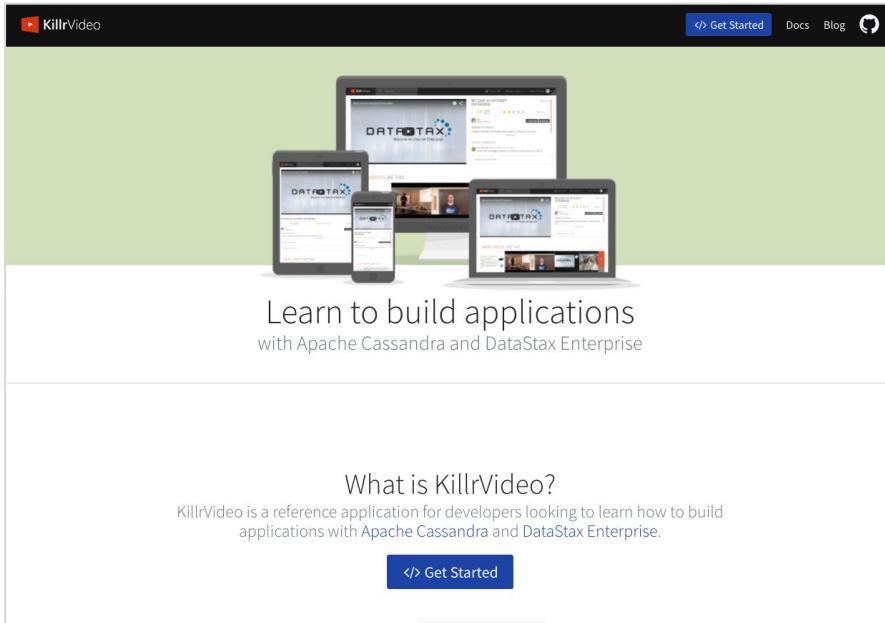
KillrVideo - a video sharing application

The screenshot displays a user interface for a video sharing application. At the top, there's a dark header bar with the text "Tour: Off" and "What is this? ▾". Below the header, the main content area features a grid of video thumbnails. Each thumbnail includes a small image, the video title, and some descriptive text below it.

- THE SECRET SCRIPTURE TRAILER #1 (2017) | MOVIECLIPS TRAILERS**
by Carmel Simonis
1 views • an hour ago
- THE GUARDIAN BROTHERS TRAILER #1 (2017) | MOVIECLIPS TRAILERS**
by Austin Dooley
1 views • an hour ago
- FERNAND TRAILER #3 (2017) | MOVIECLIPS TRAILERS**
by Jose Labadie
4 views • 2 hours ago
- THE MEYEROWITZ #1 (2017) | MOVIECLIPS TRAILERS**
by Taylor Turner
13 views • 2 hours ago

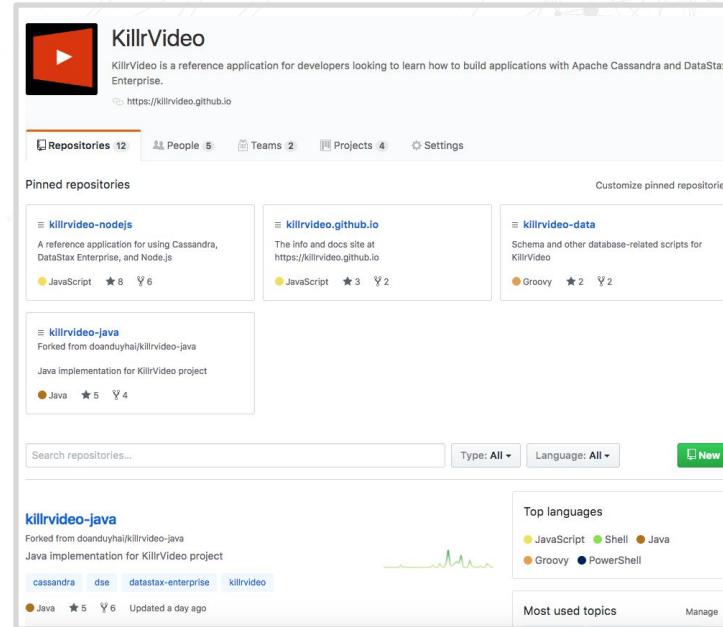
- TEST CAT VIDEOS FUNNY KITTEN SAYING**
by Arma Witting
3 views • 10 days ago
- KEY & PEELE - PEGASUS SIGHTING**
by Willie Bradtke
2264 views • 12 days ago
- HOME SWEET HELL OFFICIAL TRAILER #1 (2015) - KATHERINE HEIGL**
by Newton Reinger
2315 views • 12 days ago
- INTRODUCTION TO THE CENTER: OVERVIEW**
by Newton Reinger
2499 views • 12 days ago

KillrVideo <https://github.com/KillrVideo/killrvideo-java>



The landing page for KillrVideo features a dark header with the logo and navigation links: Get Started, Docs, Blog, and a GitHub icon. Below the header is a large image showing four devices (laptop, tablet, smartphone, and another laptop) displaying the KillrVideo application interface. A green banner below the devices contains the text: "Learn to build applications with Apache Cassandra and DataStax Enterprise". At the bottom left, there's a "What is KillrVideo?" section with a brief description and a "Get Started" button.

<https://killrvideo.github.io>

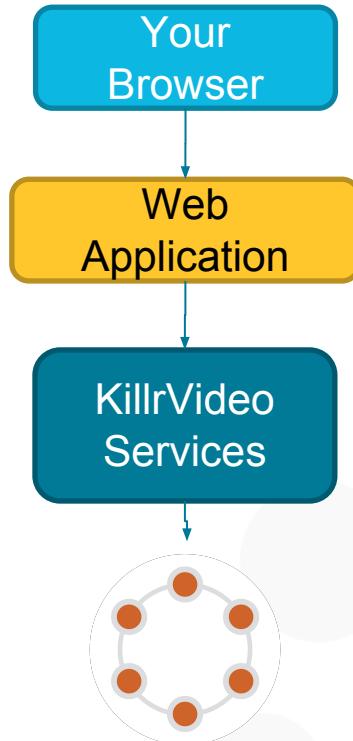


The GitHub repository page for KillrVideo shows the main repository information. It includes a video thumbnail, repository statistics (12 repositories, 5 people, 2 teams, 4 projects), and pinned repositories. One pinned repository is "killrvideo-nodejs", a reference application for using Cassandra, DataStax Enterprise, and Node.js. Another pinned repository is "killrvideo-java", a Java implementation for the KillrVideo project. The page also features a search bar, language filters, and a "New" button. On the right side, there are sections for "Top languages" (JavaScript, Shell, Java, Groovy, PowerShell) and "Most used topics".

<https://github.com/KillrVideo>



KillrVideo Architecture



Technology Choices

- Node.js
- Falcor
- Java / C# / Node.js / Python
- GRPC
- Etcd
- DataStax Drivers
- DataStax Enterprise including Apache Cassandra & Spark, Graph

Deployment

- Download and run locally via Docker
- Deployed in AWS using DataStax Managed Cloud:
<http://killrvideo.com/>

DataStax Drivers

- OSS Cassandra Drivers
 - CQL Support
 - Sync / Async API
 - Load Balancing
 - Auto Node Discovery
 - Object Mapper
- DataStax Enterprise Drivers
 - OSS Driver features plus...
 - Unified Authentication
 - Graph Fluent API
 - Geometric Types



- ODBC
- JDBC



@DataStaxAcademy #CassandraDay

Active Everywhere—Every Cloud

DATASTAX®

Maven

- Pom file (?)

```
<dependency>
  <groupId>com.datastax.dse</groupId>
  <artifactId>dse-java-driver-core</artifactId>
  <version>1.7.0</version>
</dependency>
```

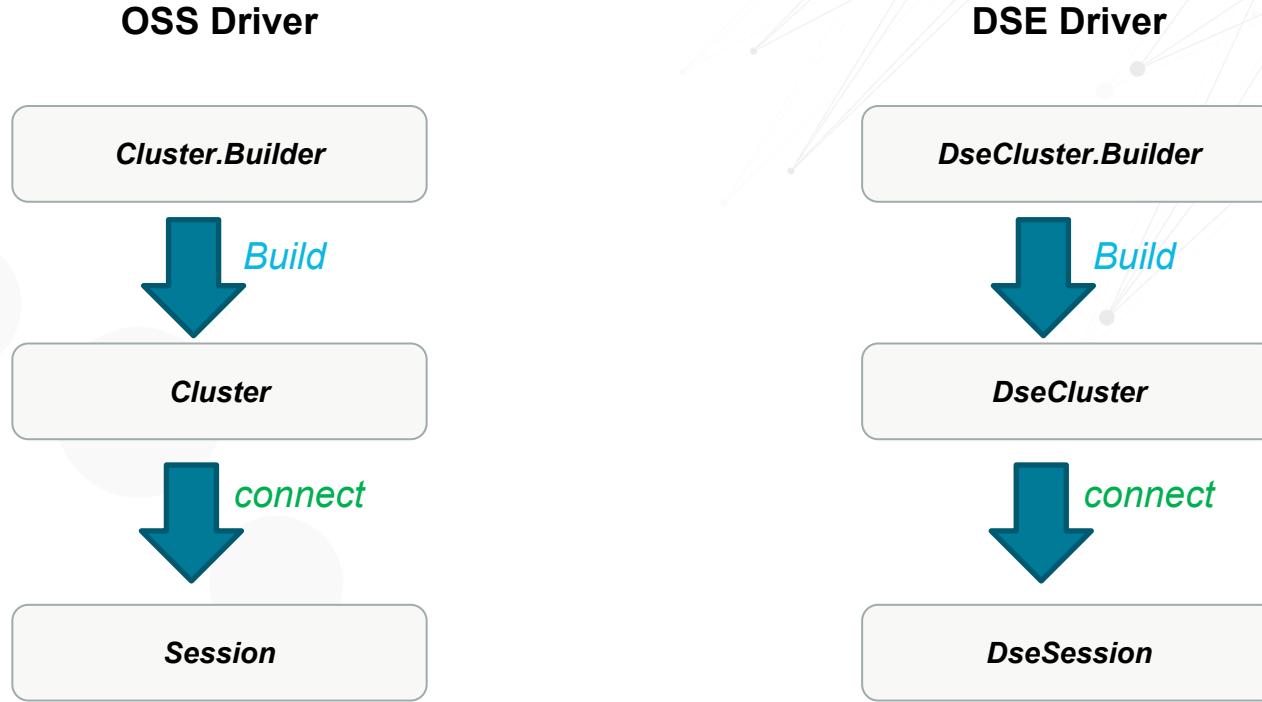
[LINK](#)



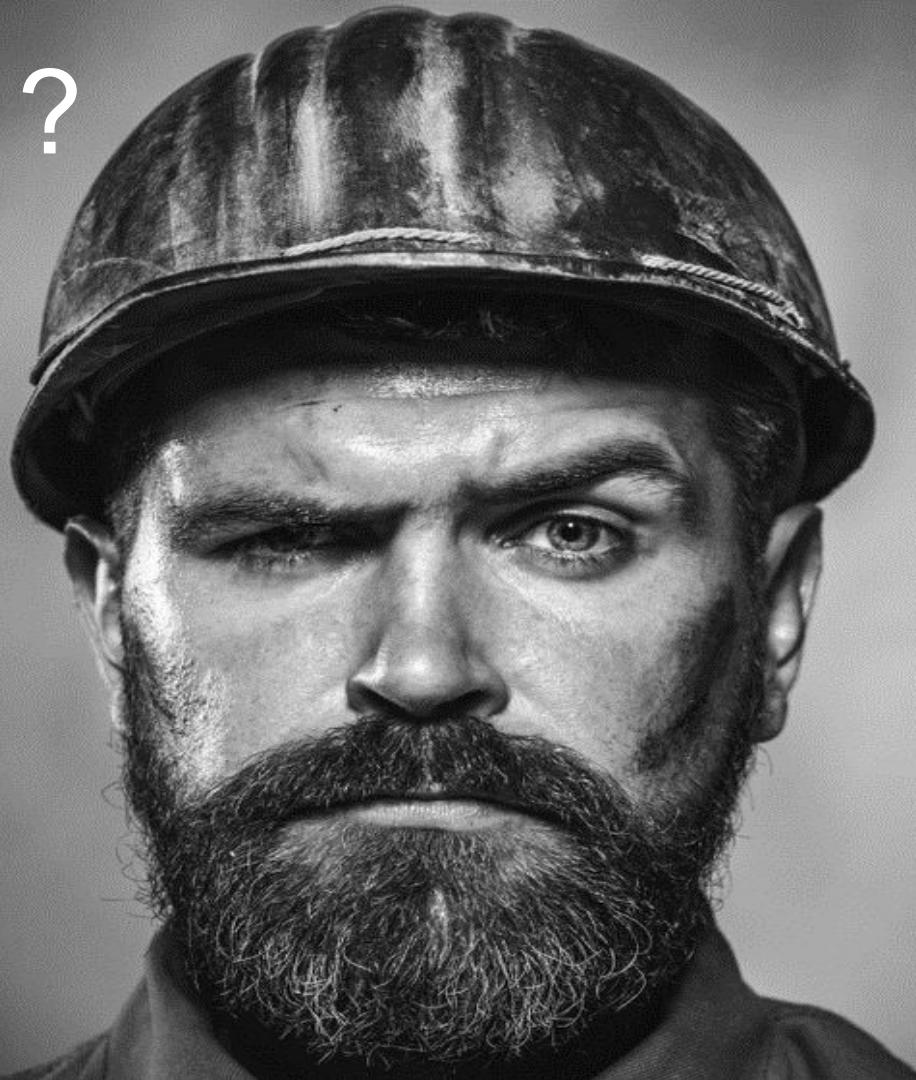
@DataStaxAcademy #CassandraDay



Connectivity



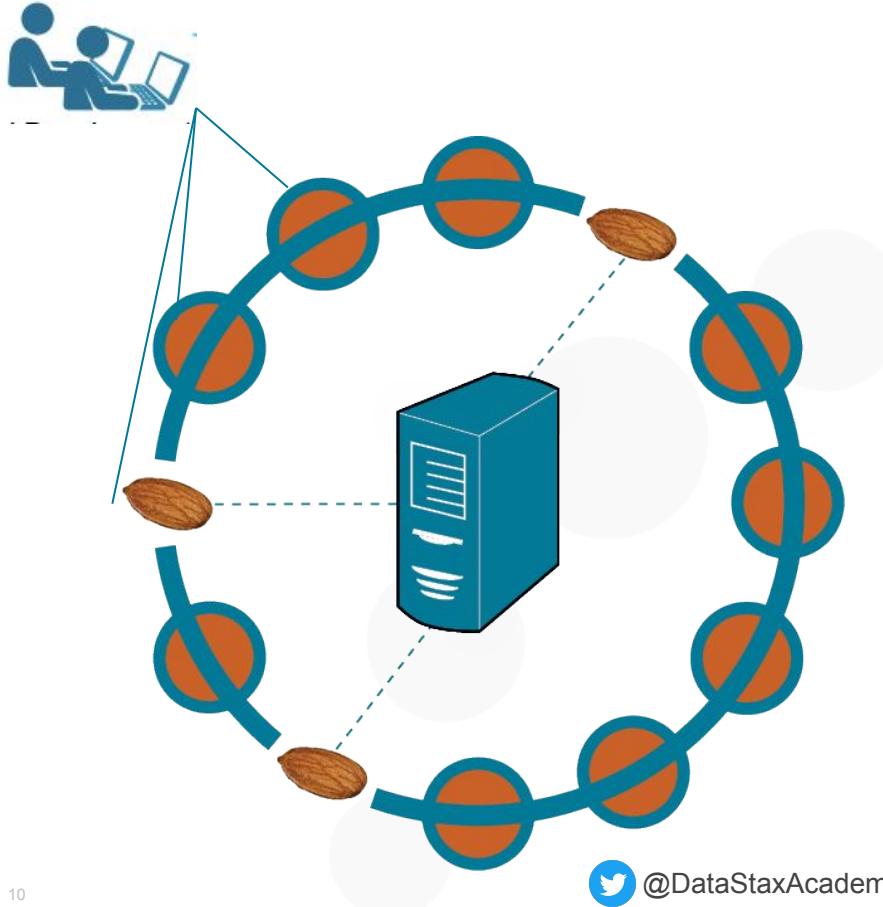
Builder ?



Builder

```
DseCluster myCluster =  
    DseCluster  
        .builder()  
        .addContactPoint("192.168.0.1")  
        .build()
```

Contact Points



Only one necessary

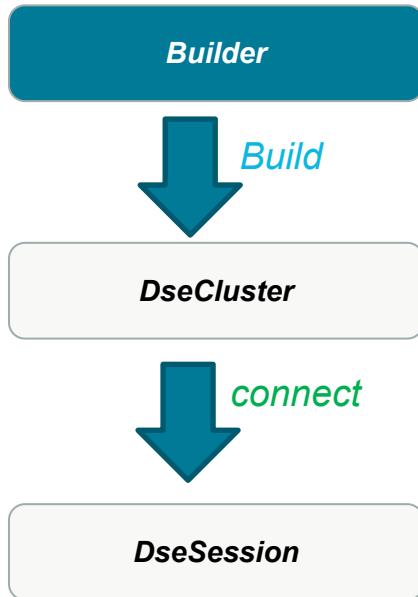
Unless that node is down



More are good

Seed nodes are good; more
“in the know”

Builder !

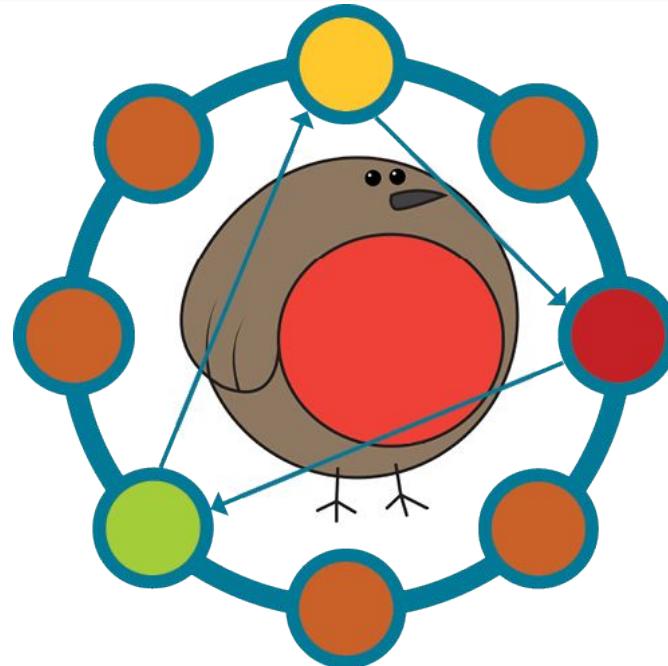


```
Builder clusterbuilder = new Builder();
clusterbuilder.addContactPoint("192.168.0.1");
clusterbuilder.addContactPoint("192.168.0.2");
clusterbuilder.addContactPoint("192.168.0.3");
clusterbuilder.withPort(9042);
```

[LINK](#)

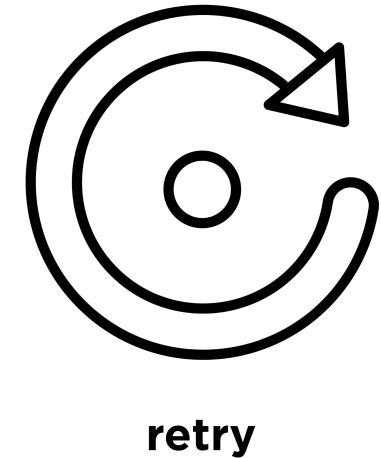
Load balancing

```
new TokenAwarePolicy(  
    new DCAwareRoundRobinPolicy("local-dc-name"))
```



Retry Policies

- Most are deprecated (long story; don't ask)
- **DefaultRetryPolicy**
 - Default
 - Retries once *onReadTimeout* or *onWriteTimeout*
 - Enough replicas for your consistency level must be online
 - Only retries idempotent mutations
- **FallthroughRetryPolicy**
 - Doesn't retry
 - Sends exception to your client application



Reconnection Policies

Reconnects driver to a downed node

Two options:

- **ConstantReconnectionPolicy**
 - Check every N milliseconds
- **ExponentialReconnectionPolicy** (default)
 - Increases every interval
 - Caps out at a max



Builder !!!!!

Builder



DseCluster



DseSession

```
Builder clusterbuilder = new Builder();
clusterbuilder.addContactPoint("192.168.0.1");
clusterbuilder.addContactPoint("192.168.0.2");
clusterbuilder.addContactPoint("192.168.0.3");
clusterbuilder.withPort(9042);

clusterbuilder.withReconnectionPolicy(
    new ExponentialReconnectionPolicy(1000, 1600));

clusterbuilder.withLoadBalancingPolicy(
    new TokenAwarePolicy(
        DCAwareRoundRobinPolicy.builder()
            .withLocalDc("DC WEST").build()));

clusterbuilder.withRetryPolicy(
    DefaultRetryPolicy.INSTANCE);

clusterConfig.withAuthProvider(
    new DsePlainTextAuthProvider("Black", "Knight"));
```

connect

Builder



DseCluster



DseSession

LINK

```
DseCluster myCluster = clusterbuilder.build();  
  
DseSession mySession = myCluster.connect("killrvideo");  
  
// DseSession mySession = myCluster.connect();
```



DseSession is all you need !

DseCluster.Builder



DseCluster



DseSession

For when you need to do something quick and dirty

```
mySession.execute("SELECT * FROM users");
```

Statement

[LINK](#)

You know...sometimes you just got to get something done



@DataStaxAcademy #CassandraDay

Active Everywhere—Every Cloud



SimpleStatement

- Driver interprets SimpleStatements at runtime and may include parameter placeholders
- You can set options on a `SimpleStatement` instance

```
Statement statement =  
    new SimpleStatement("select * from t1 where c1 = 5");  
  
Statement statement2 =  
    new SimpleStatement("select * from t1 where c1 = ?", 5);  
  
mySession.execute(statement);
```



Prepared and Bound Statements

- *Compiled once on each node automatically as needed*
- Prepare each statement only once per application
- Use one of the many bind variations to create a *BoundStatement*

```
PreparedStatement ps =  
    session.prepare("SELECT * from t1 where c1 = ?");  
  
BoundStatement bound = ps.bind(5);
```

[LINK](#)



@DataStaxAcademy #CassandraDay

Active Everywhere—Every Cloud

DATASTAX®

Query Builder



QueryBuilder

- Alternative to building CQL string queries manually
- Contains methods to build SELECT, UPDATE, INSERT and DELETE statements
- Generates a Statement as per the earlier techniques

Method

`QueryBuilder.select()`

Return Type

`Selection`

`QueryBuilder.insertInto()`

`Insert`

`QueryBuilder.update()`

`Update`

`QueryBuilder.delete()`

`Selection`

[LINK](#)



@DataStaxAcademy #CassandraDay

Active Everywhere—Every Cloud



QueryBuilder

```
QueryBuilder  
  .update("killrvideo", "video_ratings")  
  .with(QueryBuilder.incr("rating_counter"))  
  .and(QueryBuilder.incr("rating_total", QueryBuilder.bindMarker()))  
  .where(QueryBuilder.eq("videoid", QueryBuilder.bindMarker()))
```



execute() → ResultSet

```
ResultSet rs = mySession.execute("SELECT * FROM users");

// Gets all Rows (pages) at once
List<Row> myRows = rs.all();
// Better - iterate!
rs.forEach(row -> { System.out.println (row.getDouble("ColX")); });

// Working with individual Rows
Row singleRow = rs.one();
int col1 = singleRow.getInt("c1");
String col2 = singleRow.getString("c2");
```



Exercise

Cluster 1

Node 1 - 18.144.54.163

Web terminal - <https://18.144.54.163>

Spark Jobs - <http://18.144.54.163:4040/jobs>

Spark Master - <http://18.144.54.163:7080>

DataStax Studio - <http://18.144.54.163:9091>

Node 2 - 13.57.198.36

Web terminal - <https://13.57.198.36>

Eclipse Che - <http://13.57.198.36:8080/che/killrvideo>

KillrVideo - <http://13.57.198.36:3000>

Node 3 - 54.215.253.80

Web terminal - <https://54.215.253.80>

OpsCenter - <http://54.215.253.80:8888/opscenter>

DSE Application Development: Prepared Statements

Developer Day Cluster



7 days ago

...



Object Mapper

- Map Java entity beans to CQL tables

```
CREATE TABLE IF NOT EXISTS users (
    userid uuid,
    firstname text,
    lastname text,
    email text,
    created_date timestamp,
    PRIMARY KEY (userid)
);
```



```
public class User {  
  
    private UUID userid;  
    private String firstname;  
    private String lastname;  
    private String email;  
    private Date createdAt;  
  
    public User() {}  
  
    public User(UUID userid, String firstname,  
        String lastname, String email, Date createdAt) {  
        this.userid = userid;  
        this.firstname = firstname;  
        this.lastname = lastname;  
        this.email = email;  
        this.createdAt = createdAt;  
    }  
  
    // Getters and Setters  
    public UUID getUserId() { return userid; }  
    public void setUserId(UUID userid) { this.userid = userid; }  
}
```



```
@Table(keyspace = "killrvideo", name = "users")
public class User {

    @PartitionKey
    private UUID userid;

    @Column @Length(min = 1, message = "firstName must not be empty")
    private String firstname;

    @Column @Length(min = 1, message = "lastName must not be empty")
    private String lastname;

    @Column @Length(min = 1, message = "email must not be empty")
    private String email;

    @NotNull @Column(name = "created_date")
    private Date createdAt;

    // Constructors
    // Getters and Setters
}
```



Mapper

```
// Execute Once
MappingManager manager = new MappingManager(meSession);
Mapper<User> userMapper = manager.mapper(User. class);

// Usage
Result<User> users = userMapper.map(meSession.execute(meQuery));
```



Paging

- Pulling large result sets all at once may not make sense
- Retrieve sections or (pages)
- Retrieve more...if necessary
- Default of 5,000

Page Size ? `setFetchSize()`

```
// Init
Cluster cluster = Cluster.builder()
    .addContactPoint("127.0.0.1")
    .withQueryOptions(new QueryOptions().setFetchSize(2000))
    .build();

// At runtime
cluster.getConfiguration().getQueryOptions().setFetchSize(2000);

// On each
Statement statement = new SimpleStatement("your query");
statement.setFetchSize(2000);
```



[LINK](#)



@DataStaxAcademy #CassandraDay

Active Everywhere—Every Cloud

DATASTAX®

Paging State

- Tracks location in a ResultSet
- “Continue where you left off...”
- Query + parameters must match

```
ResultSet resultSet = session.execute("your query");

// iterate the result set...
PagingState pagingState =
    resultSet.getExecutionInfo().getPagingState();

// Save the state
String meState = pagingState.toString()

// Later
PagingState pagingState = PagingState.fromString(meState);
Statement st = new SimpleStatement("your query");
st.setPagingState(pagingState);
ResultSet rs = session.execute(st);
```



[LINK](#)



@DataStaxAcademy #CassandraDay

Batches and the Driver

- All commands will *eventually* succeed
- *That's it*

```
BatchStatement insertVideoBatch = new BatchStatement();  
  
insertVideoBatch.add("statement_1");  
insertVideoBatch.add("statement_2");  
insertVideoBatch.add("statement_3");  
insertVideoBatch.add("statement_4");  
  
session.execute(insertVideoBatch);
```

[LINK](#)



@DataStaxAcademy #CassandraDay

Active Everywhere—Every Cloud

DATASTAX®

Batch Details

- One message to coordinator node
- Coordinator writes a batch log and replicates it to other nodes
- All affected nodes must ack the batch before coordinator removes batch log



Batches != Data Loading

- Batches about data integrity between tables
- All three inserts must succeed for batch to succeed

```
batchStatement.add(insertVideo);  
batchStatement.add(insertUserVideo);  
batchStatement.add(insertLatestVideo);
```





**Best
Practice
Ahead**

Instantiate Less

- Single **Cluster** object per physical cluster
- Single **Session** object per application
- .. And cleanup things



Run Async

`mapper.map()` vs `mapper.mapAsync()`

`execute()` vs `executeAsync()`

`get()` vs `getAsync()`

`save()` vs `saveAsync()`

Returns a `ListenableFuture`

Maximizes Throughput

Run Async

```
public void insertComment(final Comment comment) {  
    BatchStatement batchStatement = buildBatchStatementInsertComment(comment);  
    dseSession.execute(batchStatement);  
}  
  
public CompletableFuture<Void> insertCommentAsync(final Comment comment) {  
    BatchStatement batchStatement = buildBatchStatementInsertComment(comment);  
  
    // Create a callback for processing result (hey but here it is VOID so no mapping in Success())  
    CompletableFuture<Void> cfv      = new CompletableFuture<>();  
    FutureCallback<ResultSet> myCallback = new FutureCallback<ResultSet>() {  
        public void onFailure(Throwable ex) { cfv.completeExceptionally(ex); }  
        public void onSuccess(ResultSet rs) { cfv.complete(null); }  
    };  
  
    // Bind execution and callback  
    Futures.addCallback(dseSession.executeAsync(batchStatement), myCallback);  
    return cfv;  
}
```



Prepare Once – Bind Many

- Server parses/compiles query once
- Cached for lifetime of the application
- Reduces network traffic



Avoid Deletes and Writing Nulls

```
INSERT INTO meTable (primary_key, clustering_key)  
VALUES ('pk1', 'ck1');
```

VS

```
INSERT INTO meTable (primary_key, clustering_key, regular_col)  
VALUES ('pk1', 'ck1', null);
```

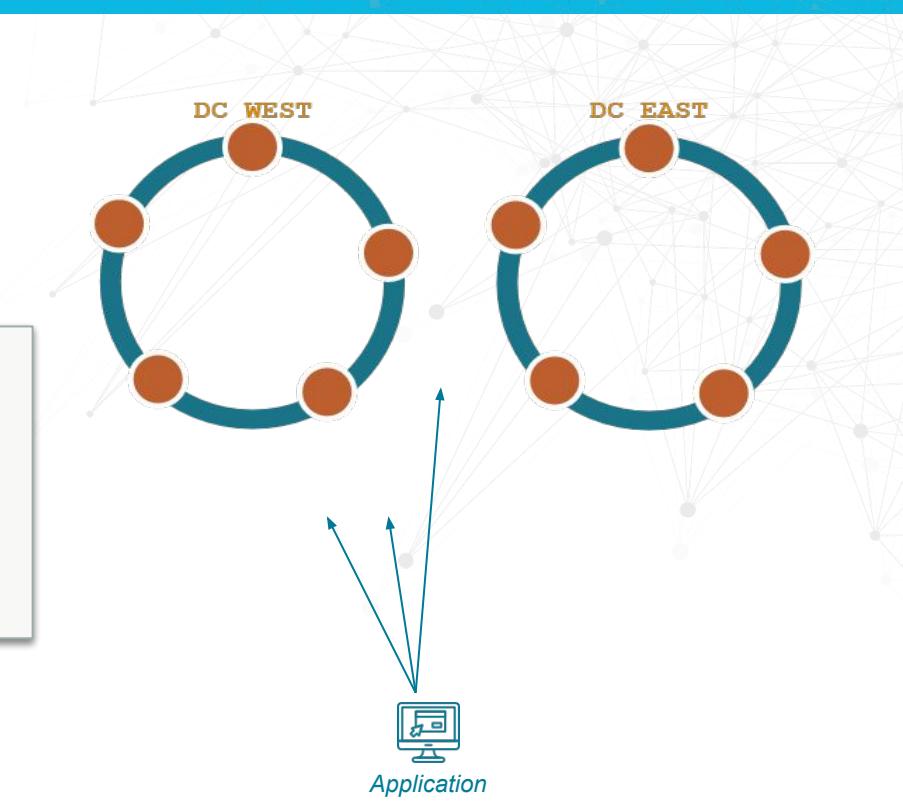
→ Second version writes a tombstone



Be Datacenter Aware

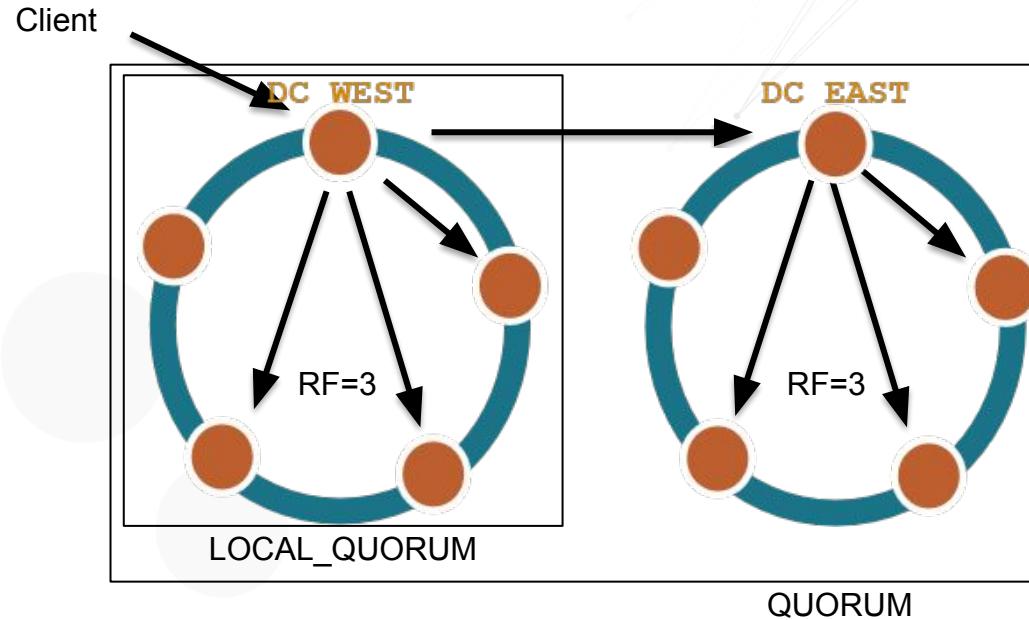
```
DCAwareRoundRobinPolicy
```

```
    .builder()  
    .withLocalDc("DC WEST")
```



Java Driver 4+ Note: You don't have to configure this

Consistency Across Data Centers



Consistency Level

```
// Global
QueryOptions options = new QueryOptions();
options.setConsistencyLevel(ConsistencyLevel.QUORUM);
Cluster cluster = Cluster
    .builder().addContactPoint("127.0.0.1")
    .withQueryOptions(options)
    .build();

// Or at runtime
cluster.getConfiguration().getQueryOptions().
setConsistencyLevel(ConsistencyLevel.QUORUM);
```

```
// On Each Query
Statement statement = new SimpleStatement("your query");
statement.setConsistencyLevel(ConsistencyLevel.QUORUM)
```





Thank you

Spring Data Cassandra



Spring Data Cassandra is generally OK to use

But if you're going to use it

Just be aware of potentially creating bad habits that could potentially crush your production database down the road



Time to Live (TTL)

- Time-series data
- Delete entire partitions when possible
- Better type of tombstones
- TimeWindowCompactionStrategy
VS DateTieredCompactionStrategy



New! DataStax Java Driver v4.0

- Targeted for Java 8 and later
- Cleans up tech debt
 - Google Guava dependency removed
- File-based configuration
 - Hot reload
 - Execution profiles
- Global timeouts
 - spanning retries, speculative executions
- Reactive APIs
- More opinionated policies
 - Single default LoadBalancingPolicy
- Object mapper in progress



<https://academy.datastax.com/content/introducing-java-driver-4>



@DataStaxAcademy #CassandraDay

Active Everywhere—Every Cloud



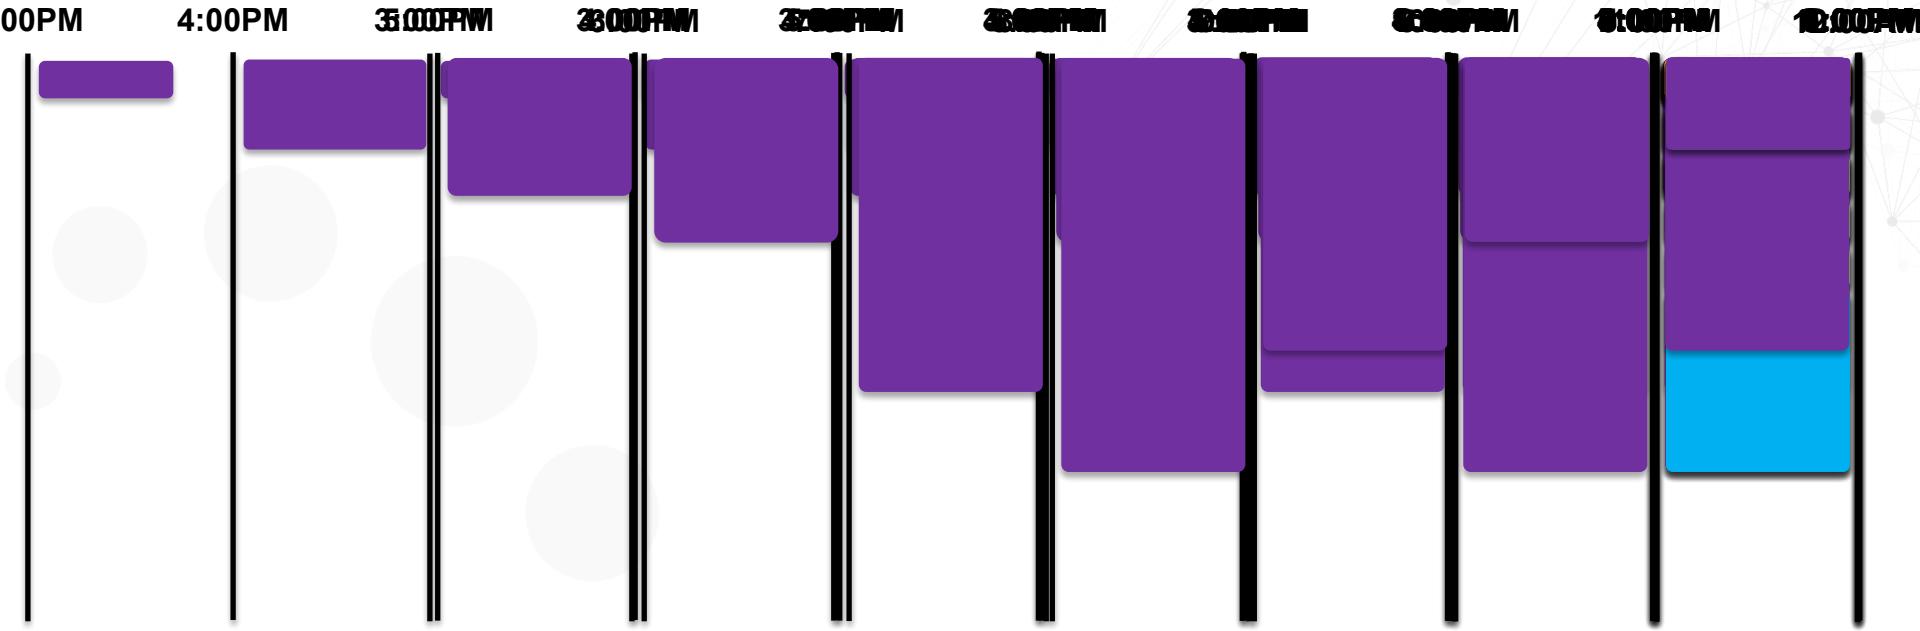
4.0 Driver API Changes - what you need to know

- Some package names changed to make a cleaner API
- Cluster combined into the Session object
 - ClusterBuilder replaced by SessionBuilder
 - Simplifies the API
- Statement objects are now immutable

https://github.com/datastax/java-driver/tree/4.x/upgrade_guide



Time Windowed Compaction

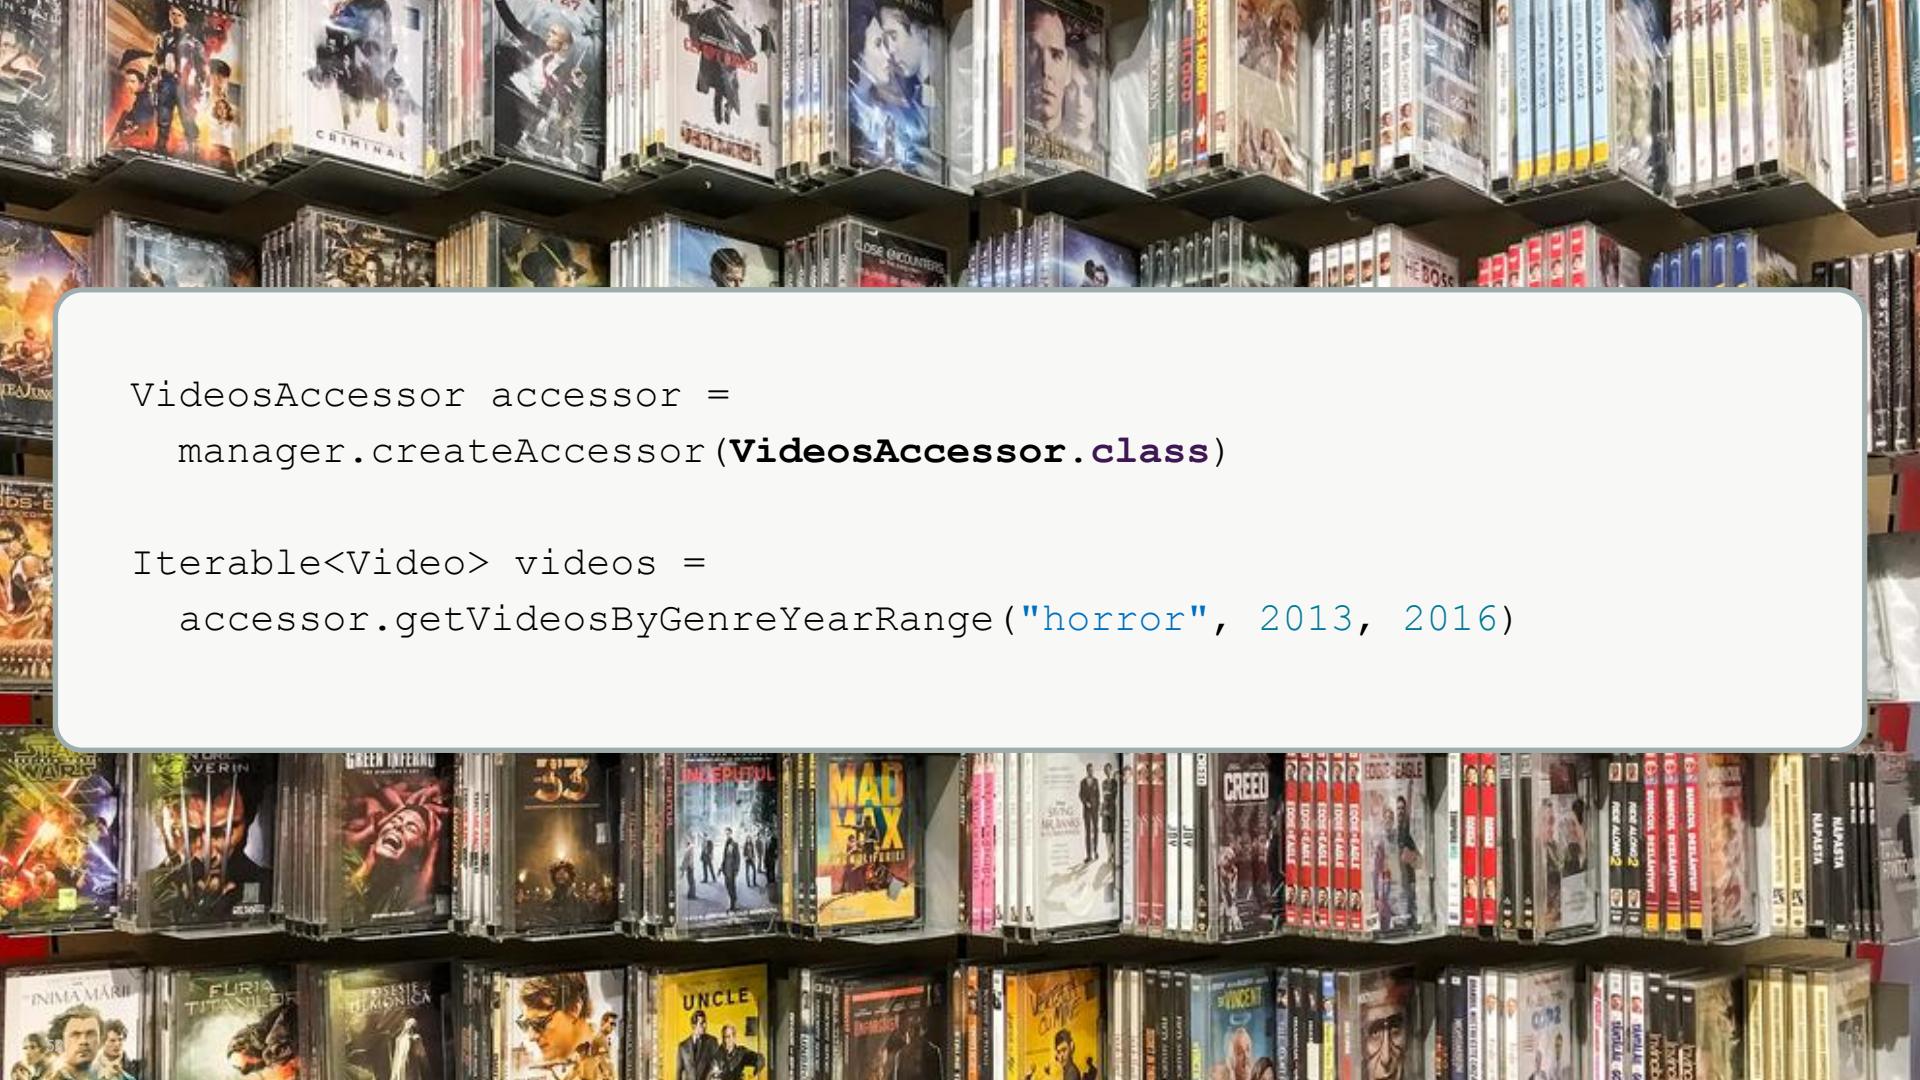


Accessors

```
@Accessor
public interface VideosAccessor {
    @Query("SELECT * FROM videos_by_genre" +
        "WHERE genre = :genre AND release_year >= :start" +
        " AND release_year <= :end ORDER BY release_year DESC ")
    Iterable<Video> getVideosByGenreYearRange(
        @Param("genre") String genre,
        @Param("start") int yearRangeStart,
        @Param("end") int yearRangeEnd);

    //...
}
```





```
VideosAccessor accessor =  
    manager.createAccessor(VideosAccessor.class)
```

```
Iterable<Video> videos =  
    accessor.getVideosByGenreYearRange("horror", 2013, 2016)
```



Exercise #2

Cluster 1

Node 1 - 18.144.54.163

Web terminal - <https://18.144.54.163>

Spark Jobs - <http://18.144.54.163:4040/jobs>

Spark Master - <http://18.144.54.163:7080>

DataStax Studio - <http://18.144.54.163:9091>

Node 2 - 13.57.198.36

Web terminal - <https://13.57.198.36>

Eclipse Che - <http://13.57.198.36:8080/che/killrvideo>

KillrVideo - <http://13.57.198.36:3000>

Node 3 - 54.215.253.80

Web terminal - <https://54.215.253.80>

OpsCenter - <http://54.215.253.80:8888/opscenter>

DSE Application Development: Accessors

Developer Day Cluster



4 days ago

...

