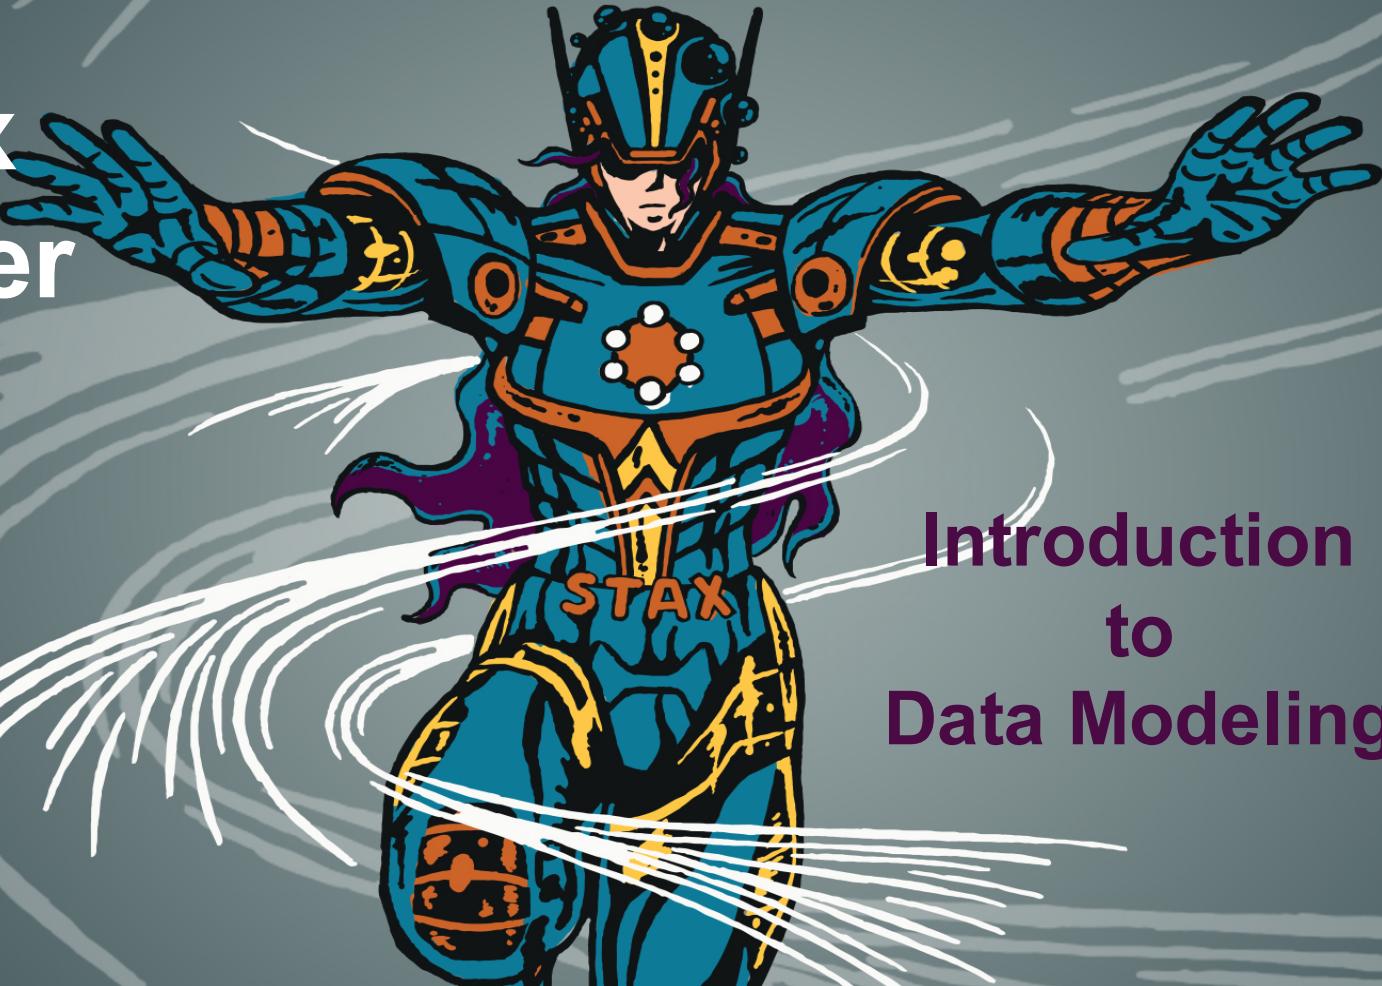


# DataStax Developer Days

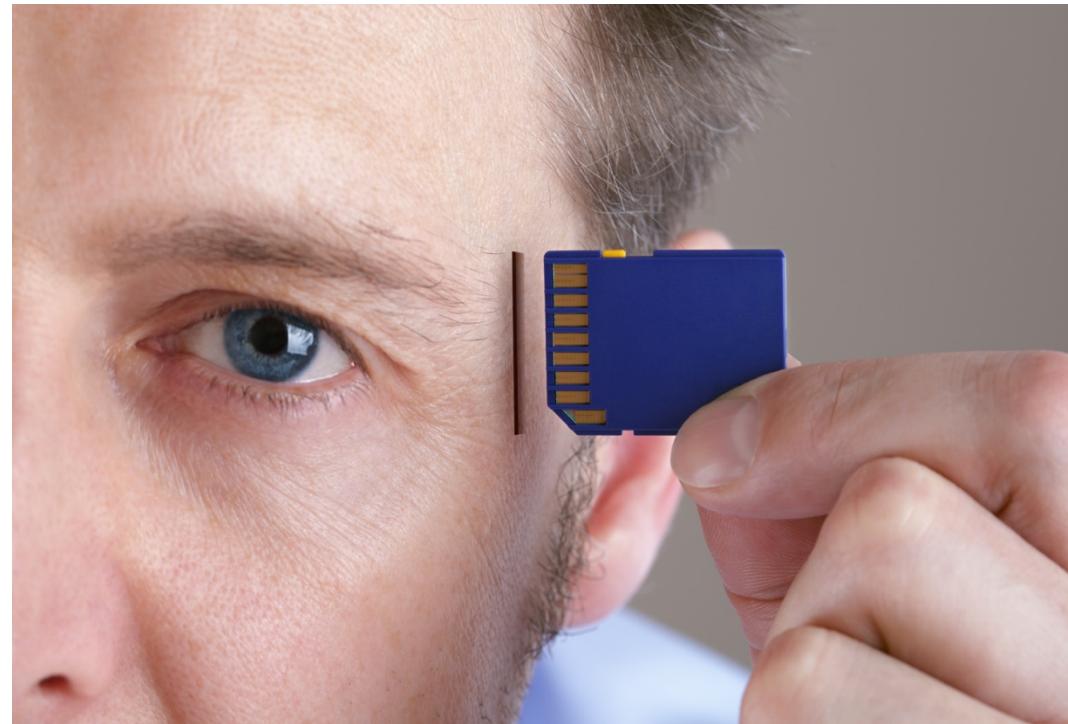


Introduction  
to  
Data Modeling

# The Cassandra Mind-shift

Change the way you think about data modeling

- Two Main concepts for successful modeling:
  - Begin with the query in mind
  - Denormalization is good



# Cassandra Data Modeling

Partition – the fundamental unit of access

- A Partition has
  - A partition key
  - Associated rows with
    - Clustering columns
    - Data columns

Partition Key:  
“Bedrock”

Hash  
Function

42

		Clustering Columns		Data Columns
		Partition		
Partition 42				
Last Name	First Name	Address	Email	
Flintstone	Dino	3 Stone St	dino@gmail.com	
Flintstone	Fred	3 Stone St	fred@gmail.com	
Flintstone	Wilma	3 Stone St	wilm@gmail.com	
Rubble	Barney	4 Rock Cir	brub@gmail.com	
Rubble	Betty	4 Rock Cir	betr@gmail.com	

# Cassandra Data Modeling

Tables hold many partitions

“Fro~~Stink~~“ Partition

Last Name	First Name	Address	Email
Moose	Bullwinkle	1 Tree St	moo@gmail.com
Squirrel	Rocky	3 Sky Blvd	fly@gmail.com
Jetson	Judy	2 StarSt	judy@gmail.com
Rubble	Barney	4 Rock Cir	brub@gmail.com
Rubble	Betty	4 Rock Cir	betr@gmail.com

Cartoon Characters by City Table

42

Last Name	First Name	Address	Email
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---

17

Last Name	First Name	Address	Email
---	---	---	---
---	---	---	---
---	---	---	---

83

Last Name	First Name	Address	Email
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---

92

Last Name	First Name	Address	Email
---	---	---	---

# Cassandra Data Modeling

Tables hold many partitions

“Bedrock” Partition

Last Name	First Name	Address	Email
Flintstone	Fred	3 Stone St	fred@gmail.com
Flintstone	Dino	3 Stone St	dino@gmail.com
Flintstone	Wilma	3 Stone St	wilm@gmail.com
Rubble	Barney	4 Rock Cir	brub@gmail.com
Rubble	Betty	4 Rock Cir	betr@gmail.com

42

Cassandra Table

Last Name	First Name	Address	Email
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---

# Cassandra Data Modeling

Tables hold many partitions

“Orbit City” Partition

Last Name	First Name	Address	Email
Jetson	George	2 Star St	gjet@gmail.com
Jetson	Jane	2 Star St	jane@gmail.com
Jetson	Judy	2 StarSt	judy@gmail.com

Cassandra Table

42

Last Name	First Name	Address	Email
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---

17

Last Name	First Name	Address	Email
---	---	---	---
---	---	---	---
---	---	---	---

# Cassandra Data Modeling

Tables hold many partitions

“Smurfville” Partition

Last Name	First Name	Address	Email
Azrael	Cat	5 Blue St	acat@gmail.com
Smurf	Brainy	7 Azure Ln	brain@gmail.com
Smurf	Clumsy	4 Teal Cir	clu@gmail.com
Smurf	Greedy	9 Royal St	gime@gmail.com
Smurf	Papa	2 Navy Ct	papa@gmail.com

Cassandra Table

42

Last Name	First Name	Address	Email
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---

17

Last Name	First Name	Address	Email
---	---	---	---
---	---	---	---
---	---	---	---

83

Last Name	First Name	Address	Email
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---

# Cassandra Data Modeling

Tables hold many partitions

“Frostbite Falls” Partition

Last Name	First Name	Address	Email
Moose	Bullwinkle	1 Tree St	moo@gmail.com
Squirrel	Rocky	3 Sky Blvd	fly@gmail.com

Cassandra Table

42	<table border="1"> <thead> <tr> <th>Last Name</th><th>First Name</th><th>Address</th><th>Email</th></tr> </thead> <tbody> <tr><td>---</td><td>---</td><td>---</td><td>---</td></tr> <tr><td>---</td><td>---</td><td>---</td><td>---</td></tr> <tr><td>---</td><td>---</td><td>---</td><td>---</td></tr> <tr><td>---</td><td>---</td><td>---</td><td>---</td></tr> <tr><td>---</td><td>---</td><td>---</td><td>---</td></tr> <tr><td>---</td><td>---</td><td>---</td><td>---</td></tr> </tbody> </table>	Last Name	First Name	Address	Email	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---
Last Name	First Name	Address	Email																										
---	---	---	---																										
---	---	---	---																										
---	---	---	---																										
---	---	---	---																										
---	---	---	---																										
---	---	---	---																										
17	<table border="1"> <tbody> <tr><td>---</td><td>---</td><td>---</td><td>---</td></tr> <tr><td>---</td><td>---</td><td>---</td><td>---</td></tr> <tr><td>---</td><td>---</td><td>---</td><td>---</td></tr> </tbody> </table>	---	---	---	---	---	---	---	---	---	---	---	---																
---	---	---	---																										
---	---	---	---																										
---	---	---	---																										
83	<table border="1"> <tbody> <tr><td>---</td><td>---</td><td>---</td><td>---</td></tr> <tr><td>---</td><td>---</td><td>---</td><td>---</td></tr> <tr><td>---</td><td>---</td><td>---</td><td>---</td></tr> <tr><td>---</td><td>---</td><td>---</td><td>---</td></tr> <tr><td>---</td><td>---</td><td>---</td><td>---</td></tr> </tbody> </table>	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---								
---	---	---	---																										
---	---	---	---																										
---	---	---	---																										
---	---	---	---																										
---	---	---	---																										
92	<table border="1"> <tbody> <tr><td>---</td><td>---</td><td>---</td><td>---</td></tr> <tr><td>---</td><td>---</td><td>---</td><td>---</td></tr> </tbody> </table>	---	---	---	---	---	---	---	---																				
---	---	---	---																										
---	---	---	---																										

# Cassandra Data Modeling

Keyspaces contain many tables

Keyspace

Cartoon Characters Table

42

Last Name	First Name	Address	Email
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---

17

---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---

83

---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---

92

---	---	---	---
---	---	---	---

Episode Table

37

Season	Episode	Name	Time
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---

47

---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---

22

---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---

38

---	---	---	---
---	---	---	---

RatingsTable

93

Cartoon	Season	Episode	Stars
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---

18

---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---

63

---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---
---	---	---	---

71

---	---	---	---
---	---	---	---

# Cassandra Data Modeling

## CQL – Cassandra's Query Language

Here's how you create a keyspace in CQL:

```
CREATE KEYSPACE cartoons
  WITH REPLICATION = {
    'class' : 'NetworkTopologyStrategy',
    'datacenter1' : 3
};
```

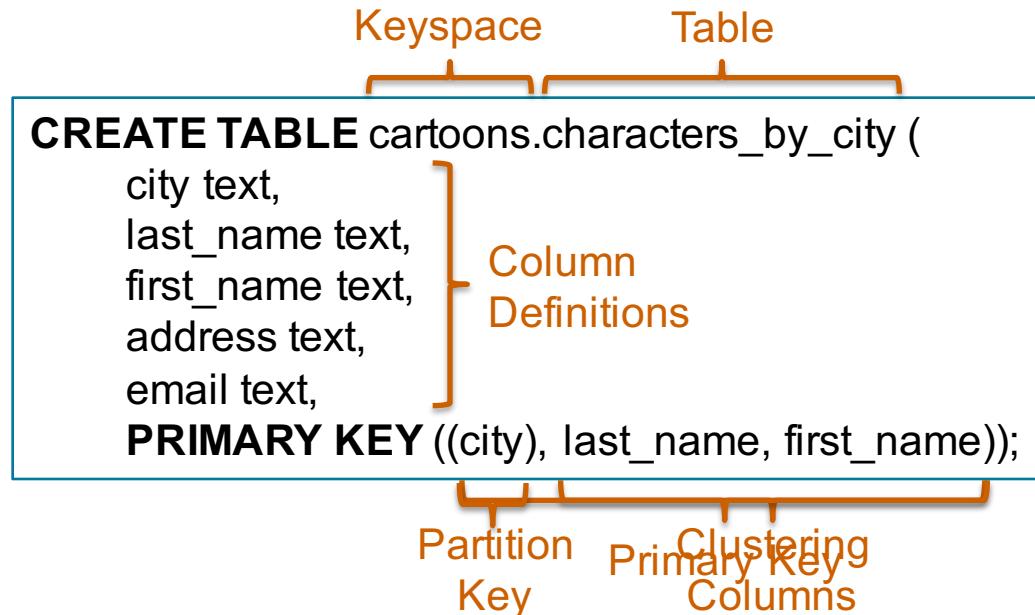
Keyspace Name

Replication Strategy

Replication Factor  
by Datacenter

# Cassandra Data Modeling

## CQL – Create table example



# Cassandra Data Modeling

## CQLSH – Investigating our schema

To list your keyspaces:

```
DESCRIBE KEYSPACES;
```

To look at a specific keyspace definition:

```
DESCRIBE KEYSPACE cartoon;
```

To look at a table definition:

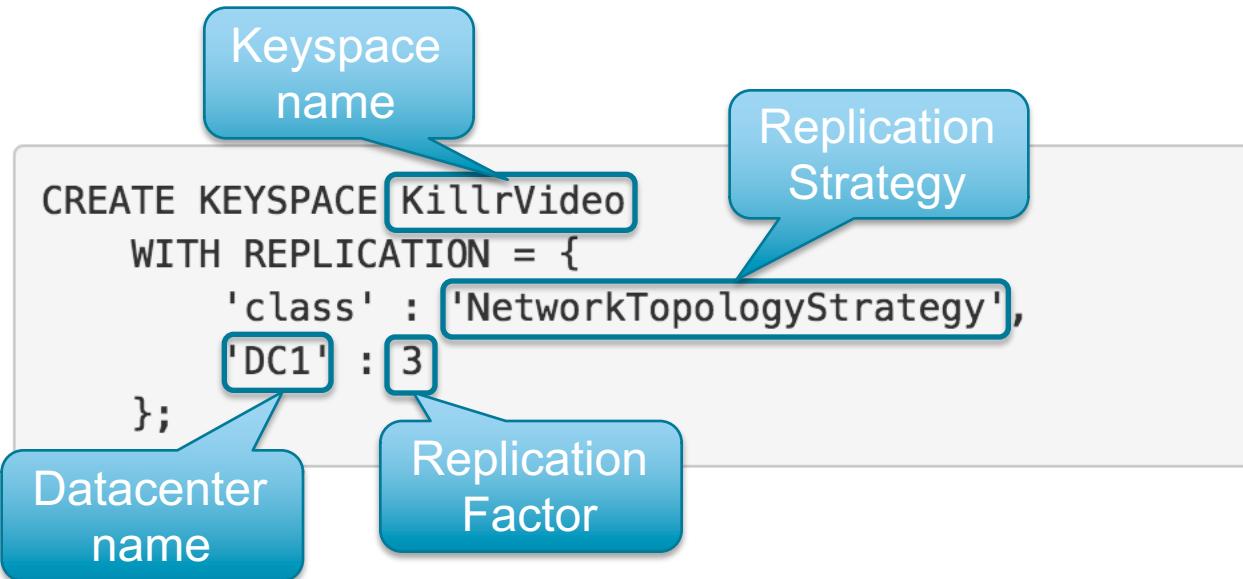
```
DESCRIBE TABLE cartoon.characters_by_city;
```

# Enough Talk Already!!!!

It's time to play!

- You each have a cluster
- You will access your cluster using DSE Studio
- Open a browser and go to <Your node's IP Address>:9091
- Click on notebook “**Data Modeling: Data Modeling Intro**”
- Work your way through the steps of this Intro





```
CREATE TABLE KillrVideo.movie_by_actor_name (
    last_name text,
    first_name text,
    movie_title text,
    movie_year int,
    PRIMARY KEY (last_name, first_name));
```

Column  
names

Table name

Column  
data types

Primary/Partition  
key

```
INSERT INTO KillrVideo.movie_by_actor_name (
    last_name,
    first_name,
    movie_title,
    movie_year)
VALUES(
    'Stone',
    'Emma',
    'Battle of the Sexes',
    2017);
```

Column  
order

Column  
data

```
CREATE TABLE KillrVideo.movie_by_actor_name (
    last_name text,
    first_name text,
    movie_title text,
    movie_year int,
    PRIMARY KEY ((last_name, first_name), movie_title));
```

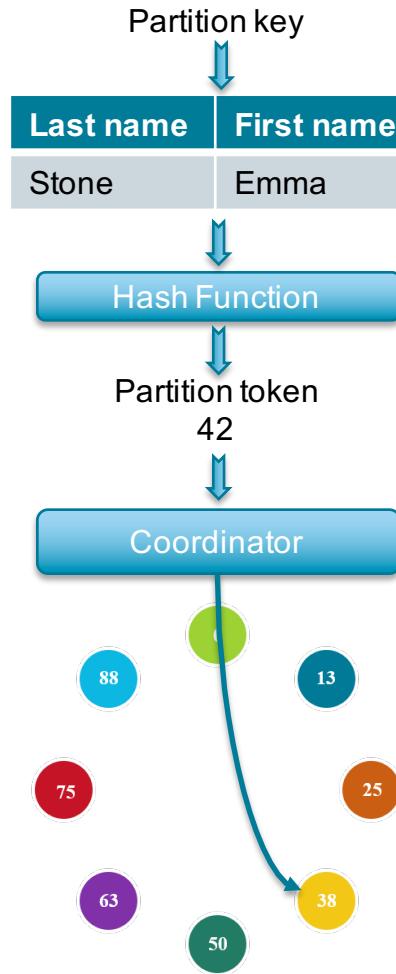
Entire primary key

Partition key

Clustering Column

Last name	First name	Movie title	Movie year
Depp	Johnny	Lone Ranger	2013
Depp	Johnny	Tourist	2010
Jolie	Angelina	Tourist	2010
Jolie	Angelina	By the Sea	2015
Pitt	Brad	Big Short	2015
Pitt	Brad	Moneyball	2011





# Data Modeling Intro

## Quick review

- Key concepts:
  - Keyspace – contains tables
  - Table – contains partitions
  - Partition – basic unit of storage/retrieval
    - Identified by partition key embedded within primary key
    - Contains one or more rows
  - Row – has a primary key and data columns
  - Primary key – intra-table row identifier
    - Consists of partition key and clustering columns
  - Partition key – partition identifier, hashes to partition token
  - Clustering column – intra-partition key for sorting rows within partition



# Let's do a Data Modeling Project Together!

# Welcome to Cassandra-Land

The Theme Park Where You Can Find...

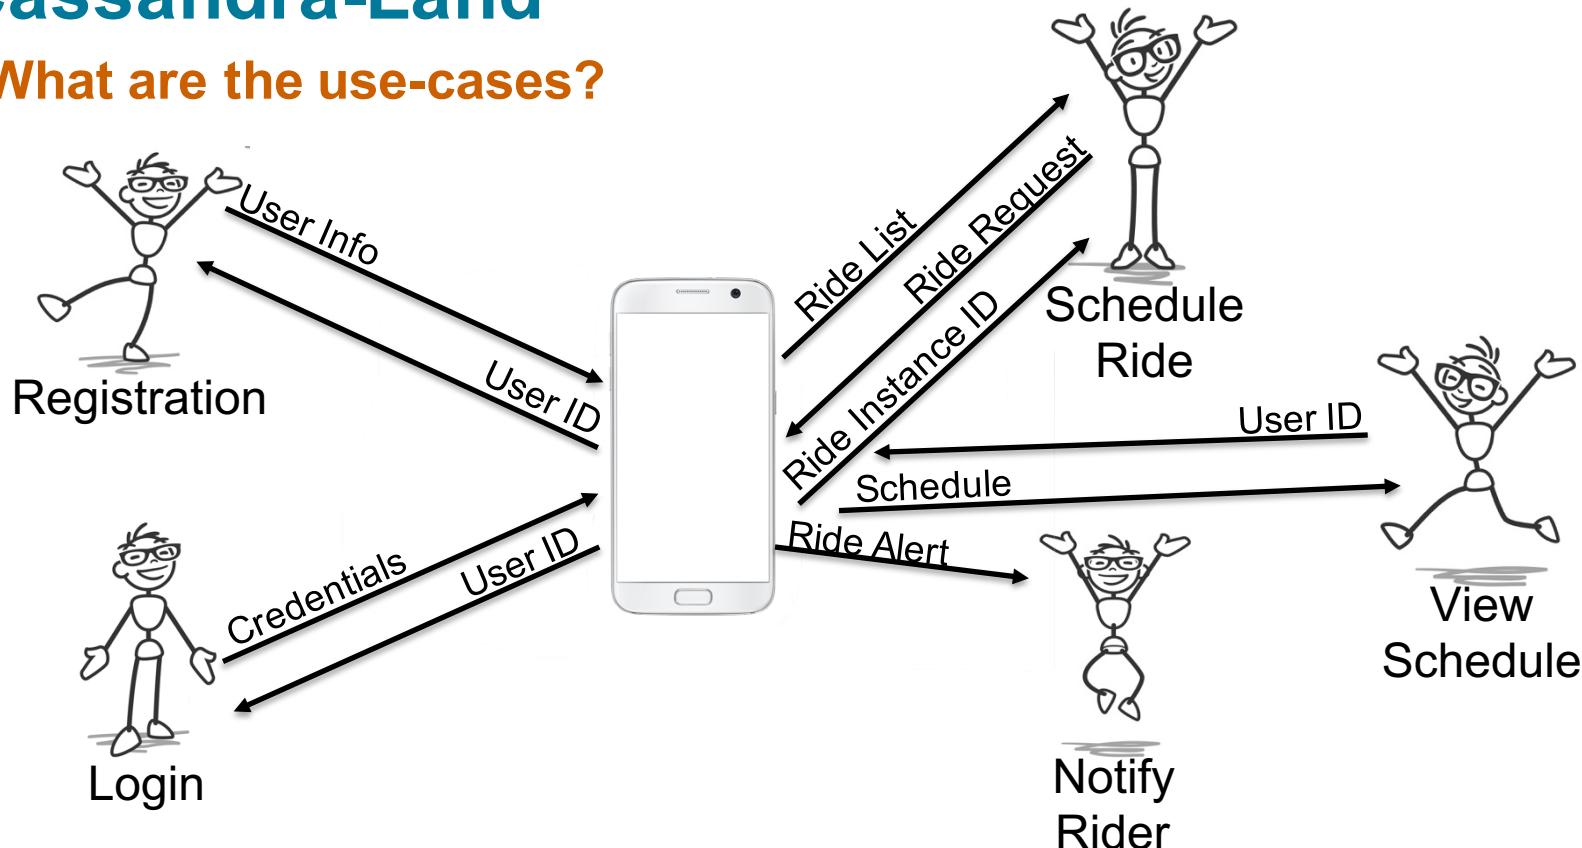
- Distributed & Fault-Tolerant Rides
- Amazing Throughput
- And Fast Response Times

But We Need an App!



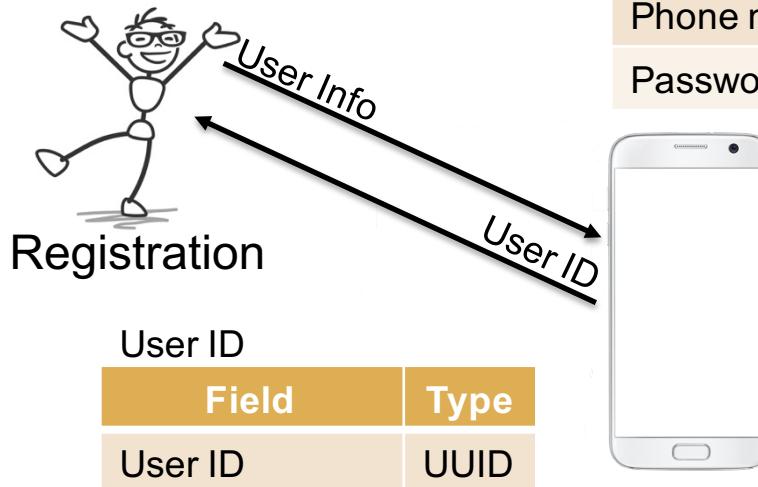
# Cassandra-Land

## What are the use-cases?



# Cassandra-Land

## Registration Use-Case



users\_by\_phone\_number

Field	Type
phone_number	text
user_id	UUID
password	text

# Registration Use-Case

## What you need to know – CREATE KEYSPACE

```
CREATE KEYSPACE <keyspace name> WITH REPLICATION = {  
    'class' : <replication strategy>,  
    <datacenter name> : <replication factor>,  
    // Specify addition datacenters/replication factors here...  
};
```

### For example:

```
CREATE KEYSPACE killrvideo WITH REPLICATION = {  
    'class' : NetworkTopologyStrategy,  
    'USWestDC': 3,  
    'USEastDC': 2  
};
```

# Registration Use-Case

## What you need to know – DESCRIBE KEYSPACE

```
DESCRIBE KEYSPACE <keyspace name>;
```

# Registration Use-Case

## What you need to know – CREATE TABLE

```
CREATE TABLE <keyspace name>.<table name> (
    <field name> <field type>,
    // Add additional field descriptions here
    PRIMARY KEY(<primary key descriptor>)
);
```

# Registration Use-Case

## What you need to know – INSERT

```
INSERT INTO <keyspace name>.<table name>
(<column list>)
VALUES(<column values>);
```

# Registration Use-Case

What you need to know – **SELECT \***

```
SELECT * FROM <keyspace name>.<table name>;
```

# Registration Use-Case

## What you need to know – Upserts

- Cassandra does **NOT** read before writing
- Inserting a row with the same primary key causes an update
- This behavior is called an “upsert”
- Similarly, updates to non-existent rows cause an insert
- **INSERT INTO keyspace.table IF NOT EXISTS ...**
  - Lightweight transaction to prevent upserts

# Notebook Data Modeling: Cassandra-Land Project

## Do Steps 1A and 1B

# Cassandra-Land

## Login Use-Case

Credentials

Field	Type
Phone number	text
Password	text



Login

Credentials

User ID



User ID

Field	Type
User ID	UUID

users\_by\_phone\_number

Field	Type
phone_number	text
user_id	UUID
password	text

# Registration Use-Case

## What you need to know – SELECT

```
SELECT * FROM <keyspace name>.<table name>  
    WHERE <query constraints>;
```

- Must include full partition key
- Partition keys do NOT support inequalities
- Not all clustering columns need be specified, but...
- Any preceding clustering columns *MUST* be specified

# Notebook Data Modeling: Cassandra-Land Project

**Do Step 2**

# Cassandra-Land

## Ride Alert Use-Case

ride\_instances\_by\_start\_time

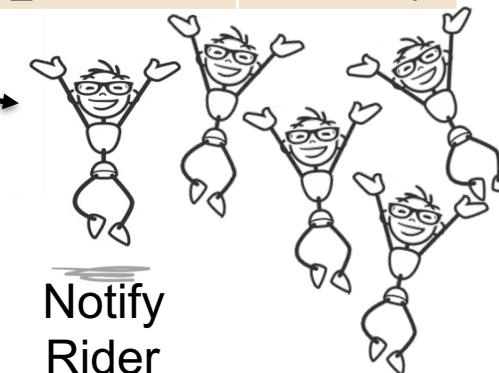
Field	Type
<b>start_time</b>	timestamp
<i>ride_id</i>	UUID↑
ride_name	text
<i>user_id</i>	UUID↑
phone_number	text



Ride Alert

Field	Type
phone_number	text
ride_name	text
<b>start_time</b>	timestamp

Ride Alert



# Cassandra-Land

## View Schedule Use-Case

ride\_instances\_by\_user\_id

Field	Type
<b>user_id</b>	UUID
<i>start_time</i>	timestamp↑
<b>ride_id</b>	UUID
<b>ride_name</b>	text



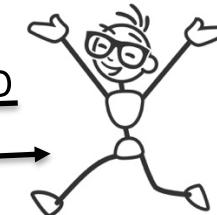
User ID

Field	Type
User ID	UUID

Schedule

Schedule

Field	Type
<i>start_time</i>	timestamp
<b>ride_name</b>	text



View  
Schedule

# Cassandra-Land

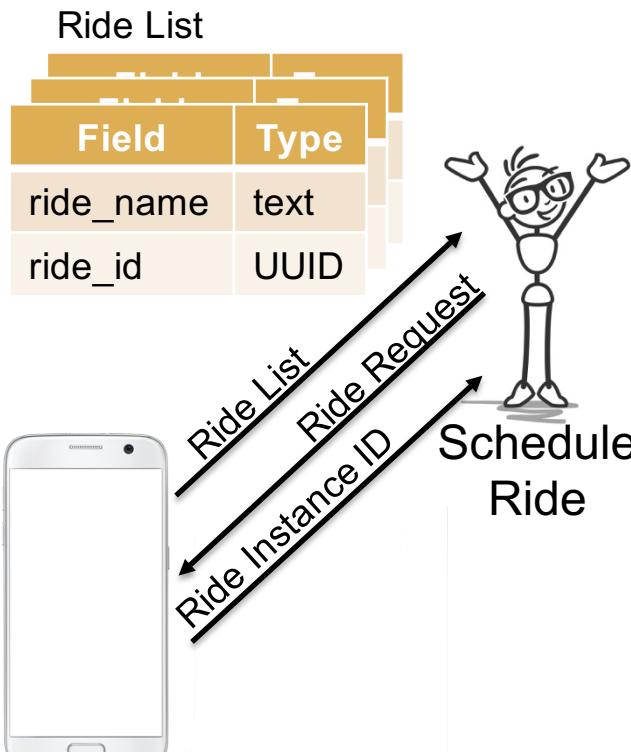
## Schedule Ride Use-Case

ride\_list\_by\_location

Field	Type
location	text
ride_id	UUID↑
ride_name	text
capacitiy	int

rider\_count\_by\_time\_and\_ride

Field	Type
start_time	timestamp
ride_id	UUID
rider_count	int



DATASTAX  
A C A D E M Y

Ride Request

Field	Type
user_id	UUID
ride_id	UUID
start_time	timestamp

Ride Instance ID

Field	Type
ride_instance_id	UUID

# Cassandra-Land

## Table Summary

users\_by\_phone\_number

Field	Type
phone_number	text
user_id	UUID
password	text

ride\_list\_by\_location

Field	Type
location	text
ride_id	UUID↑
ride_name	text
capactiy	int

rider\_count\_by\_time\_and\_ride

Field	Type
start_time	timestamp
ride_id	UUID
rider_count	int

ride\_instances\_by\_user\_id

Field	Type
user_id	UUID
start_time	timestamp↑
ride_id	UUID
ride_name	text

ride\_instances\_by\_start\_time

Field	Type
start_time	timestamp
ride_id	UUID↑
ride_name	text
user_id	UUID↑
phone_number	text

# Last Three Use-Cases

## What you need to know – PRIMARY KEY()

PRIMARY KEY(( <partition key column>, ...), <clustering column>, ...)

- Must have one or more partition key columns
- May have zero or more clustering columns

# Last Three Use-Cases

## What you need to know – Timestamps

Format is 'YYYY-MM-DDTHH:MM:SS'

- Notice the quotes

# Last Three Use-Cases

## What you need to know – UPDATE

UPDATE <keyspace name>.<table name>

    SET <assignment>

    WHERE <row specification>

IF <condition>

- Can have multiple <assignment>
- IF is optional – causes a lightweight transaction

# Last Three Use-Cases

## What you need to know – BATCH

BEGIN BATCH

  INSERT statement

  INSERT statement

  ...

APPLY BATCH

- Causes all operations to complete
- Use for inserting into multiple tables

# Notebook Data Modeling: Cassandra-Land Project

## Do Steps 3-5

Phone Number	Password
8017124927	JamiesPhone-NotReally
9999999999	NinesAreWild
1234567890	SeeQuential

# Cassandra-Land

## What you learned!

- How to analyze use-cases to derive a data model
- How to denormalize to maintain performance
- How to use lightweight transactions
- How to leverage batch operations





The End